
Distill~~Coder~~Math: Capturing Large Model Performance in a Small Model via Distillation

Group 7 - Zhiyang Zhang, Jake Watson

Why bother distilling?

Reasoning over memorization

Supervised Fine-Tuning teaches the model to memorize patterns by minimizing Cross-Entropy loss.

Distillation teaches the model to learn the reasoning of the teacher by adding a KL-Divergence term to the loss function.


Where we started

Model	Performance on GSM8K	VRAM	Toks/s
Qwen2.5-Math-7B-Instruct	85.4%	14GB	~50
Qwen2.5-1.5B-Instruct	60-70%	3GB	~250

Table 4: Performance of the 7B+ base models.

Datasets	Mistral-7B	Llama3-8B	Gemma2-9B	Qwen2-7B	Qwen2.5-7B
<i>Mathematics & Science Tasks</i>					
GPQA	24.7	25.8	32.8	30.8	36.4
TheoremQA	19.2	22.1	28.9	29.6	36.0
MATH	10.2	20.5	37.7	43.5	49.8
MMLU-stem	50.1	55.3	65.1	64.2	72.3
GSM8K	36.2	55.3	70.7	80.2	85.4

The Process



Pipeline

Collect

Teacher generates
answers to math
problems → JSON
dataset


Distill

Student trains to
match teacher's
logits + ground-
truth answer

Eval

Test student on
GSM8K

Stage 1: Collection




Prompt Template:

```
<|im_start|>system
Please reason step by step, and put
your answer in \boxed{}.
<|im_end|><|im_start|>user
{QUESTION}<|im_end|>
<|im_start|>assistant
```


Output:

```
{prompt, [response_1, response_2,
...]}
```


Stage 2: Distillation




Training log



GSM8K - 1316 test examples - pass@1 - 8GB max VRAM



GSM8K - 1316 test examples - pass@1 - 8GB max VRAM



GSM8K - 1316 test examples - pass@1 - 8GB max VRAM

Metric	Teacher (Qwen-2.5-Math-7B-Instruct)	Student (Qwen-2.5-1.5B-Instruct)	Distilled (Ours)
Latency	-	16.9ms/tok	17.1ms/tok
Peak VRAM	OOM	3.31	3.29GB
Energy Avg (GPU) nvidiaDeviceGetPowerUsage	-	114 W	113 W
Energy / correct	-	197 J	158 J
Accuracy	85.4%	66.3%	80.5%

"question": "A plague infects ten people. Every day, each infected person infects six others. How many people are infected after three days?",

"gold": "3430",

"responses": {

"Distilled 1.5B": {

"full": "To determine how many people are infected after three days, we need to follow the spread of the plague step by step.\n\n1. ****Initial Condition:****\n - On day 0, there are 10 infected people.\n\n2. ****Day 1:****\n - Each of the 10 infected people infects 6 others.\n - Number of new infections on day 1: $(10 \times 6 = 60)$ \n - Total number of infected people on day 1: $(10 + 60 = 70)$ \n\n3. ****Day 2:****\n - Each of the 70 infected people infects 6 others.\n - Number of new infections on day 2: $(70 \times 6 = 420)$ \n - Total number of infected people on day 2: $(70 + 420 = 490)$ \n\n4. ****Day 3:****\n - Each of the 490 infected people infects 6 others.\n - Number of new infections on day 3: $(490 \times 6 = 2940)$ \n - Total number of infected people on day 3: $(490 + 2940 = 3430)$ \n\nThus, the total number of people infected after three days is $(\boxed{3430})$."

"predicted": "3430",

"correct": true

},

"Qwen 1.5B Instruct": {

"full": "To determine how many people are infected after three days, we need to follow the process of infection day by day.\n\n****Day 0:****\n- Initially, there is 1 infected person (the original plague).\n\n****Day 1:****\n- Each of the 1 infected person infects 6 new people.\n- Therefore, on Day 1, $(1 \times 6 = 6)$ new infections occur.\n- Total number of infected people after Day 1: $(1 + 6 = 7)$ \n\n****Day 2:****\n- Each of the 7 infected people from Day 1 infects 6 new people.\n- Therefore, on Day 2, $(7 \times 6 = 42)$ new infections occur.\n- Total number of infected people after Day 2: $(7 + 42 = 49)$ \n\n****Day 3:****\n- Each of the 49 infected people from Day 2 infects 6 new people.\n- Therefore, on Day 3, $(49 \times 6 = 294)$ new infections occur.\n- Total number of infected people after Day 3: $(49 + 294 = 343)$ \n\nThus, the total number of people infected after three days is $(\boxed{343})$."

"predicted": "343",

ours

base

Demo (Video)

<https://youtu.be/ks5di1iFiXQ>